

# Enhanced Image Metadata Extractor

## Project Members (Group B)

- Merlin Volkmer
- Bence Balázs
- Florian Johann Hafner

## Language

Python & Terraform (for deployment)

## Provider

Google Cloud

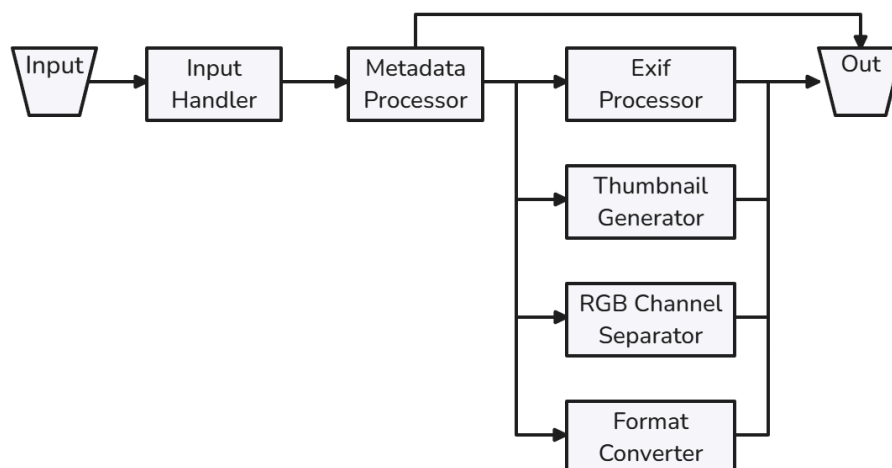
## Input

An image file (png, jpg, jpeg)

## Output

JSON metadata and converted images

## Workflow Design



## Connections between the functions

Each function is independent to the other, but it's execution is triggered by the previous step using Pub/Sub (i.e the execution of the Format Conversion is triggered by the Pub/Sub topic "convert-image")

# Functions

## Input Handler (`input_handler/main.py`)

**Trigger:** Cloud Storage event (object finalized in input bucket)

**Purpose:** Initial entry point for processing pipeline

**Inputs:**

- Cloud Event with file metadata
- File formats: JPG/JPEG/PNG only

**Outputs:**

- PubSub messages to `extract-metadata` topic
- Error response for invalid formats

**Process:**

1. Validates file extension
2. Publishes messages to metadata extraction topic
3. Hardcodes output bucket to `serverless-project-24w-output`

## Metadata Extractor (`metadata_extractor/main.py`)

**Trigger:** PubSub message from `extract-metadata` topic

**Purpose:** Extract basic image metadata

**Inputs:**

- PubSub message with file location
- Image from input bucket

**Outputs:**

- JSON metadata file in output bucket (`metadata/{filename}.json`)
- PubSub messages to 4 processing topics:
  - `generate-thumbnail`
  - `convert-format`
  - `separate-rgb-channels`
  - `process-exif`

**Collected Metadata:**

- Format, dimensions, color mode, timestamp

## EXIF Processor (`exif_processor/main.py`)

**Trigger:** PubSub message from `process-exif` topic

**Purpose:** Extract EXIF metadata

**Inputs:**

- Image from input bucket
- Existing metadata JSON

**Outputs:**

- Updated metadata JSON with EXIF data
- Handles binary EXIF data conversion

**Special Handling:**

- Converts binary data to strings
- Error tracking for problematic tags

## Format Converter (`format_converter/main.py`)

**Trigger:** PubSub message from `convert-format` topic

**Purpose:** Convert image to multiple formats

**Inputs:**

- Source image from input bucket

**Outputs:**

- Converted images in 3 formats: JPEG, PNG, WEBP
- Metadata JSON with conversion paths
- Files stored in `converted/{filename}/` path

**Note:** Skips conversion to original format

## RGB Channel Separator (`rgb_channel_separator/main.py`)

**Trigger:** PubSub message from `separate-rgb-channels` topic

**Purpose:** Split image into RGB components

**Inputs:**

- Source image from input bucket

**Outputs:**

- 3 channel images (red, green, blue) as PNGs
- Metadata JSON with channel paths
- Files stored in `channels/{filename}/` path

**Preprocessing:**

- Converts non-RGB images to RGB mode

## Thumbnail Generator (`thumbnail_generator/main.py`)

**Trigger:** PubSub message from `generate-thumbnail` topic

**Purpose:** Create resized thumbnails

**Inputs:**

- Source image from input bucket

**Outputs:**

- Thumbnails in 2 sizes: 200x200 and 400x400
- Maintains original image format
- Metadata JSON with thumbnail paths
- Files stored in `thumbnails/{filename}/` path

## Testing

### Load Test

The load Test axis features two modes that differ in the time spent to send images and the amount of concurrent processes running for that duration.

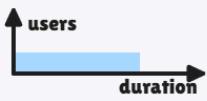







- simple
  - 5 concurrent users uploading images for 10 seconds
- production
  - baseline: 5 concurrent users uploading images for 60 seconds
  - medium: 10 concurrent users uploading images for 100 seconds
  - stress: 100 concurrent users uploading images for 3 seconds

### Terraform configuration

The Terraform configuration axis features three modes that differ by the available memory allocated to each function.

- Low: 128 **MiB (Mebibytes)**
- Mid: 256 **MiB**
- High: 512 **MiB**

This results in the following six scenarios.

Scenario	Memory		
	low	128Mi	mid 256Mi high 512Mi
<b>Simple</b> 	<b>Scenario 1 (S1)</b> 	<b>Scenario 2 (S2)</b> 	<b>Scenario 3 (S3)</b> 
<b>Production</b> 	<b>Scenario 4 (S4)</b> 	<b>Scenario 5 (S5)</b> 	<b>Scenario 6 (S6)</b> 

## Results

All data collected during the tests can be found in the results directory of the submitted archive.

### Observations


- `rgb-channel-seperator` and `format-converter` need significantly more resources than the other functions
- functions crash due to being out of memory if they are allocated less than 128 Mi

### Recommendations

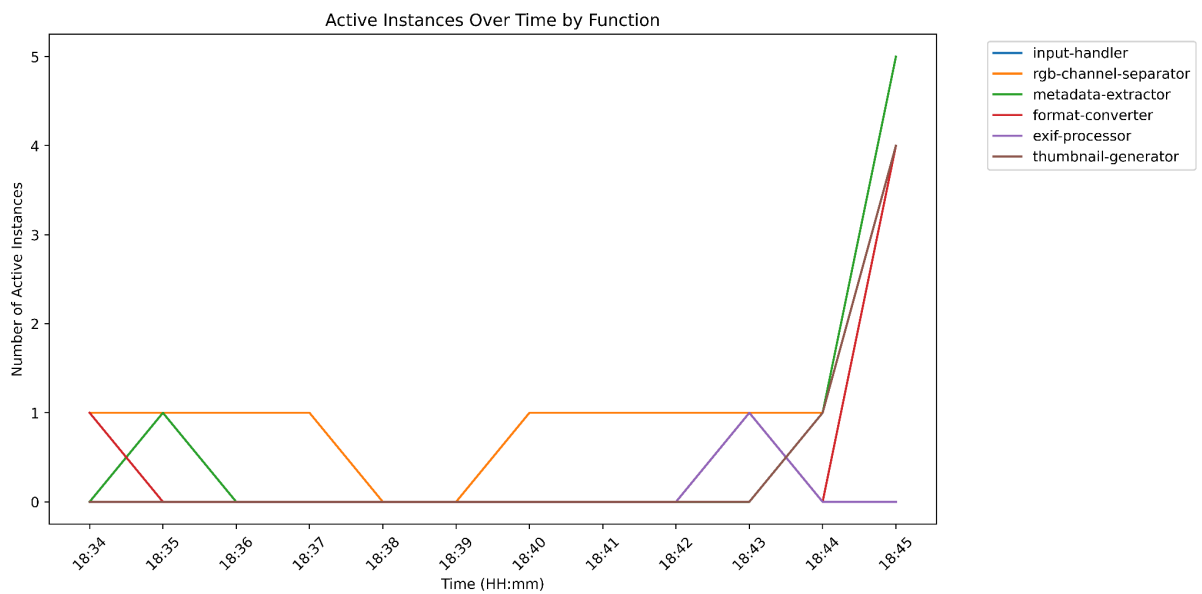
- all functions operate on images and thus need at least 128 Mi
- `rgb-channel-seperator` and `format-converter` should be allocated more resources than other functions

# Active Instances

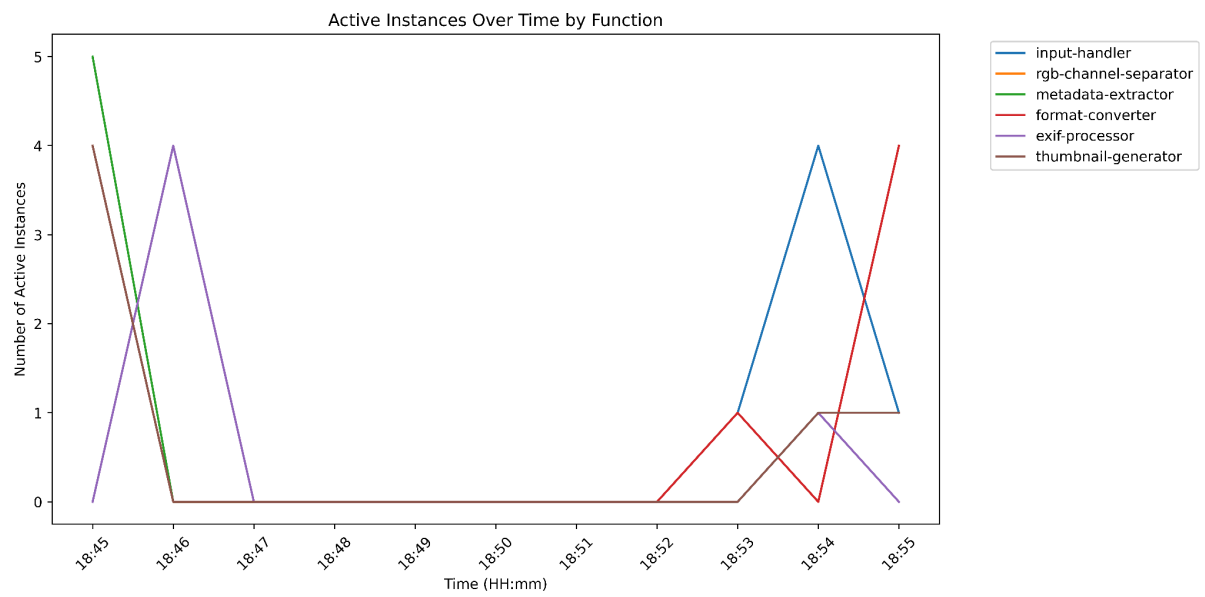
R1




R2



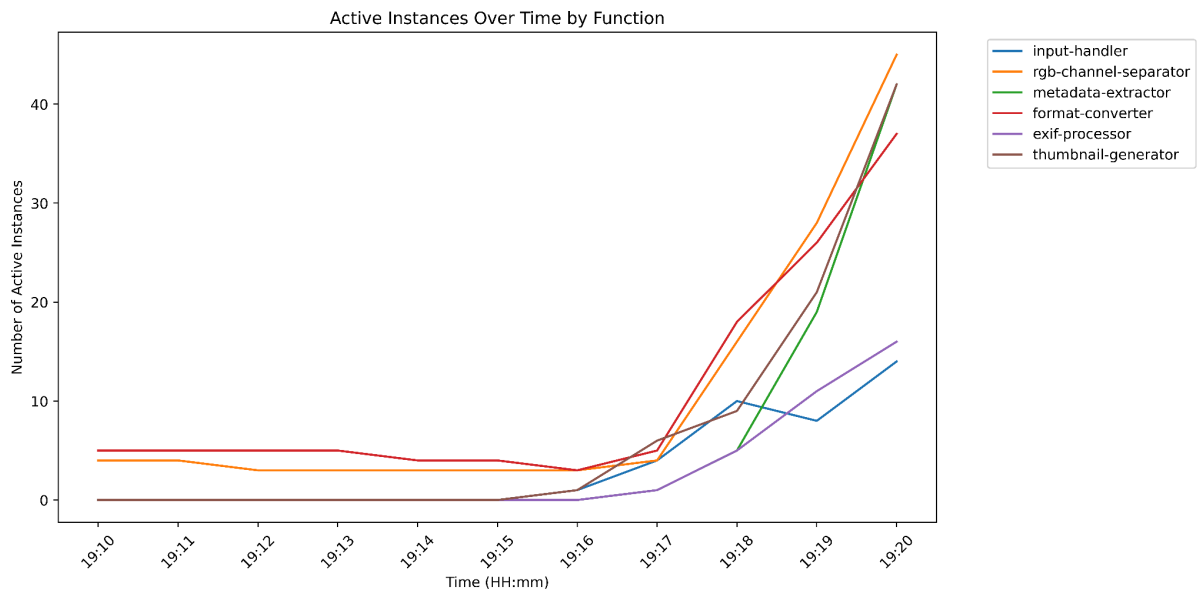
R3



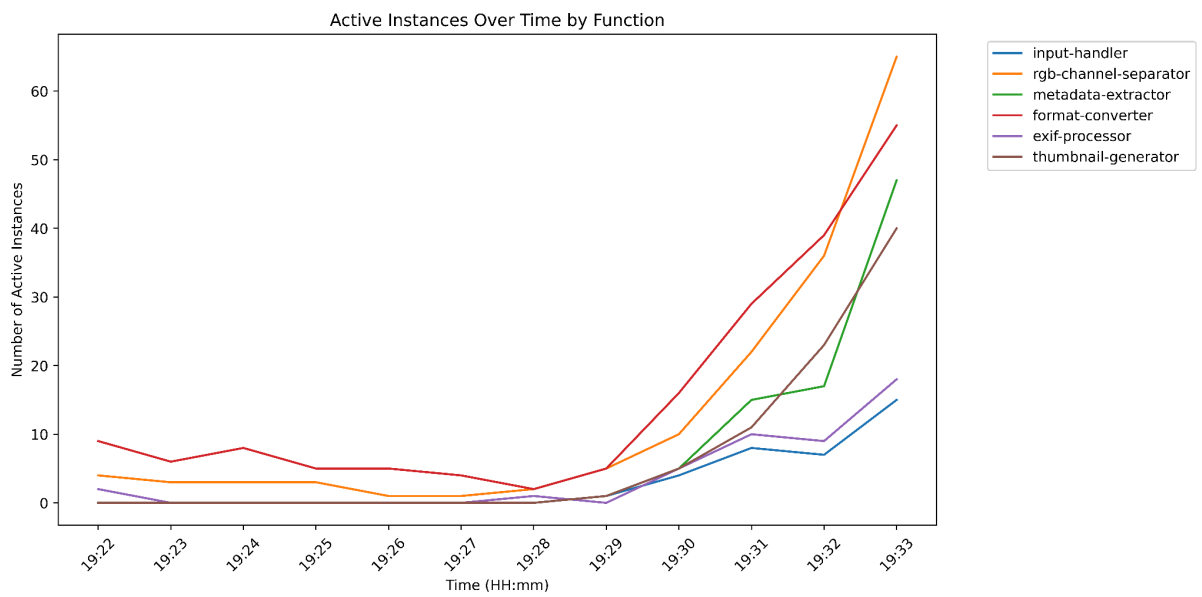
R4



R5




R6






# Memory Consumed


R1




R2




R3




R4



R5



R6



## Challenges, Considerations and Closing Words

The primary challenge encountered during the project solely related to GCPs IAM configurations, specifically in establishing the necessary permissions of the cloud user accounts for communication between the different triggers, pub/sub topics, storage events and cloud functions.

Despite these early setbacks, the project provided valuable insights for all of us into the advantages and challenges associated with serverless architectures and cloud-based computing.